

COMP 6231:


Distributed System Design



Introduction to Distributed Systems

Based on Chapters 1 and 2 of the textbook and slides from
Professor Wolfgang Emmerich, London University
Professor Kenneth P. Birman, Cornell University

What is a Distributed System?





What is a Distributed System?

- A distributed system is a collection of autonomous hosts that are connected through a computer network. Each host executes components and operates a distribution middleware, which enables the components to coordinate their activities in such a way that *users perceive the system as a single, integrated computing facility.*



How to hide Distribution of Components?

- A components should be able to communicate with any other component the same way.
 - *Access Transparency*
- Components should be able to communicate using their names only without knowing their location
 - *Location Transparency*



Transparency

- Distributed systems should be perceived by users and application programmers as a whole rather than as a collection of cooperating components.
- Transparency has different dimensions that were identified by ANSA.
- These represent various properties that distributed systems should have.



Transparency in a Distributed System

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource may be replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource
Persistence	Hide whether a (software) resource is in memory or on disk



Common Requirements

- Certain requirements are common to many distributed systems
 - Concurrency
 - Resource Sharing
 - Openness
 - Scalability
 - Fault Tolerance
 - Transparency



Concurrency

- Components in distributed systems are executed in concurrent processes.
- Components access and update shared resources (e.g. variables, databases, device drivers).
- Integrity of the system may be violated if concurrent updates are not coordinated.
 - Lost updates
 - Inconsistent analysis



Resource Sharing

- Ability to use any hardware, software or data anywhere in the system.
- Resource manager controls access, provides naming scheme and controls concurrency.
- Resource sharing model (e.g. client/server or object-based) describing how
 - resources are provided,
 - they are used and
 - provider and user interact with each other.



Openness

- Openness is concerned with extensions and improvements of distributed systems.
- Detailed interfaces of components need to be published.
- New components have to be integrated with existing components.
- Differences in data representation of interface types on different processors (of different vendors) have to be resolved.



Scalability


- Adaption of distributed systems to
 - accommodate more users
 - respond faster (this is the hard one)
- Usually done by adding more and/or faster processors.
- Components should not need to be changed when scale of a system increases.
- Design components to be scalable!



Fault Tolerance


- Hardware, software and networks fail!
- Distributed systems must maintain availability even at low levels of hardware/software/network reliability.
- Fault tolerance is achieved by
 - recovery
 - redundancy

Failure Tolerance



- There is at any time a single primary process and one or more secondary (backup, slave) processes
- FEs communicate with the primary which executes the operation and sends copies of the updated data to the result to backups
- If the primary fails, one of the backups is promoted to act as the primary

High Availability



- All the processes perform the same computation
- A single result is produced from the results of all the processes
- Correct result can be produced even when one of the processes produces incorrect result



Why Distributed Systems – What problems are solved?

- Easily connect users to **remote resources**
- Share resources with remote users in a controlled way
 - Hide the fact that the resources are physically distributed over a network -- **transparency**
 - Should be an **open** system
 - Offers services by standard rules that describe the syntax and semantics of those services
 - Should be **scalable**
 - Size, geography, and administration



Distributed System Characteristics

- Multiple autonomous components
- Components are not shared by all users
- Resources may not be accessible
- Software runs in concurrent processes on different processors
- Multiple Points of control
- Multiple Points of failure



Centralized System Characteristics

- One component with non-autonomous parts
- Component shared by users all the time
- All resources accessible
- Software runs in a single process
- Single Point of control
- Single Point of failure



Examples of mission-critical distributed applications

- Banking, stock markets, stock brokerages
- Health care, hospital automation
- Control of power plants, electric grid
- Telecommunications infrastructure
- Electronic commerce and electronic cash on the Web (very important emerging area)
- Corporate “information” base: a company’s memory of decisions, technologies, strategy
- Military command, control, intelligence systems



We depend on distributed systems!

- If these critical systems don't work
 - When we need them
 - Correctly
 - Fast enough
 - Securely and privately
- ... then revenue, health and safety, and national security may be at risk!
- *Should be scalable, failure-tolerant and highly available*



Major Issues in Distributed Systems Design

- **Independent Failure**: Nodes can fail, rest should go on
- **Unreliable and Insecure Communication**: Connections may fail, messages may be contaminated or modified
- **Costly Communication**: Communication cost (latency etc) is higher than local communication within a node
- **High Availability, Fault-Tolerance** are key advantages
- **Atomicity is difficult** due to absence of Global Time

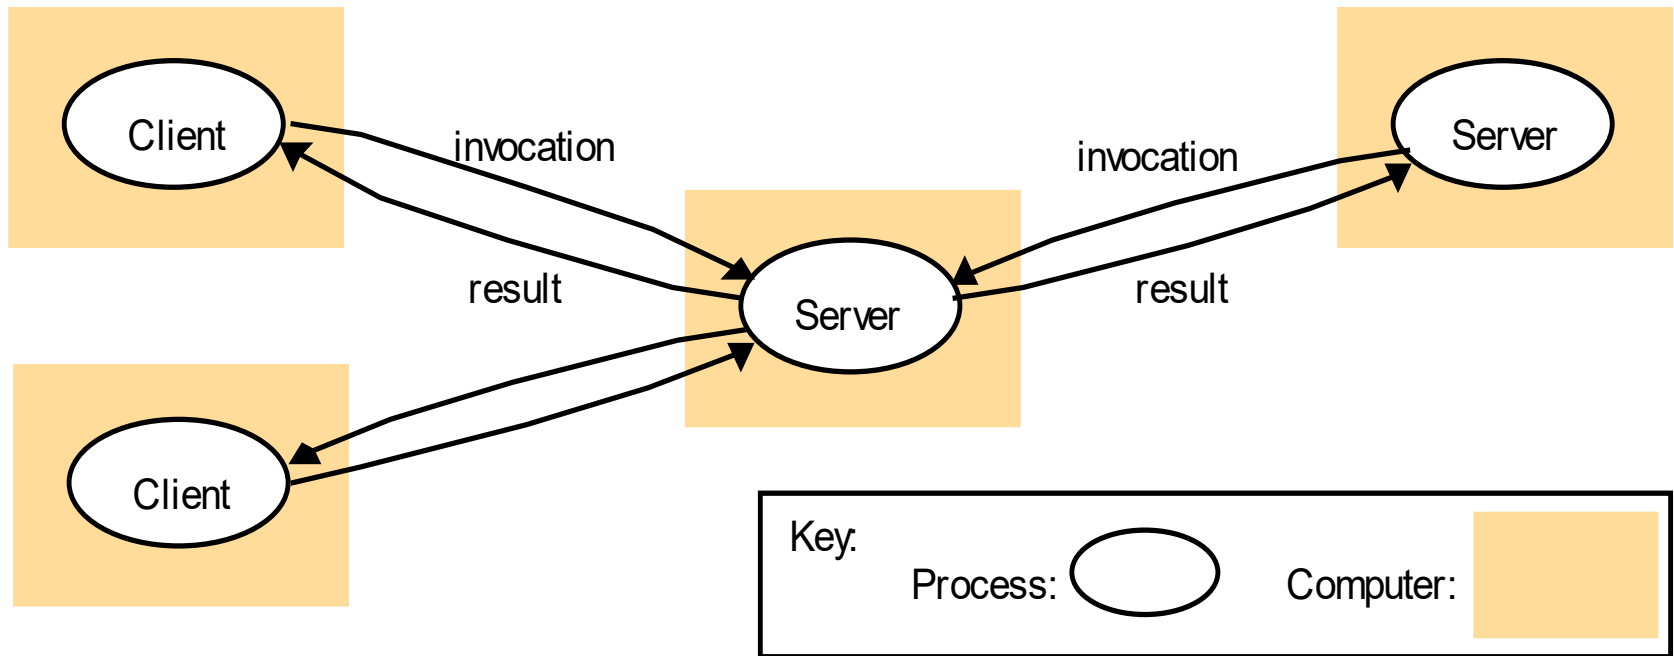


Distributed Computing Paradigms

- There are many distributed computing paradigms/architectures
 - Client-Server,
 - Peer-to-Peer,
 - Distributed Agents,
 - Distributed Object Spaces, ...
- We will focus on the client-server paradigm only.

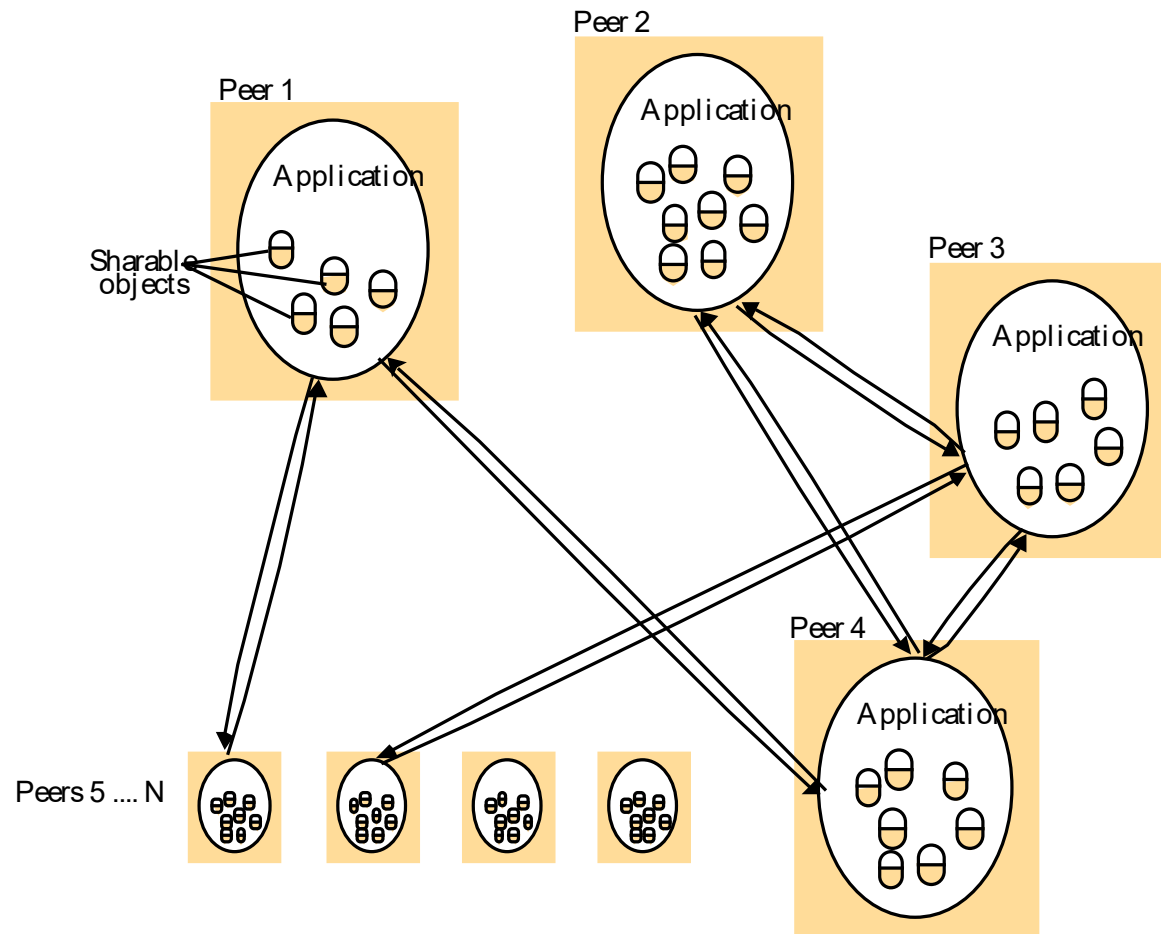
Distributed Systems

Architecture: Client-Server



Distributed Systems

Architecture: Peer-to-Peer





Client-Server Interaction

- Client access to server should be independent of
 - Local or remote server
 - Access Transparency
 - Location of the server
 - Location Transparency
 - Implementation of the server
 - Implementation Transparency,
 - Programming Language Transparency,
 - Replication Transparency, ...



Remote Procedure Call (RPC)

- Pre-RPC: Most applications were built directly over the Internet primitives
 - Their idea: mask distributed computing system using a “transparent” abstraction
 - Looks like normal procedure call
 - Hides all aspects of distributed interaction
 - Supports an easy programming model
- Today, RPC is the core of many distributed systems
- Today, RPC often used from object-oriented systems employing CORBA or COM standards.



RPC is fundamental

- **RMI**: Object Oriented RPC
- **Network Services**: RPC among mostly homogeneous processes (over LANs)
- **CORBA**: RPC among heterogeneous processes (over WANs using IIOP)
- **Web Services**: RPC over the web using HTTP



CORBA

- The Common Object Request Broker Architecture (CORBA) is a standard architecture for a distributed objects system.
- CORBA is designed to allow distributed objects to interoperate in a heterogenous environment, where objects can be implemented in different programming language and/or deployed on different platforms



Web Services

- Today, we normally use Web browsers to talk to Web sites
 - Browser names document via URL (lots of fun and games can happen here)
 - Request and reply encoded in HTML, using HTTP to issue request to the site
- Web Services generalize this model so that computers can talk to computers



Server Design

- Many of the server characteristics are achieved through process replication
 - Scalable
 - Automatically create server processes to distribute client requests
 - Fault Tolerant
 - Maintain multiple server processes so that a faulty server can be taken over by a good one
 - Highly Available
 - Maintain multiple server processes so that a client request can be handled without delay